

Predicting a song is hit or not, based on its the audio features.

Importing Libraries

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
```

In [2]:

```
df = pd.read_csv('dataset.csv')
```

In [3]:

```
df.shape
```

Out[3]:

```
(114000, 21)
```

In [4]:

```
df.columns[df.isnull().any()]
```

Out[4]:

```
Index(['artists', 'album_name', 'track_name'], dtype='object')
```

In [5]:

```
df.isnull().sum()
```

Out[5]:

```
Unnamed: 0      0
track_id        0
artists         1
album_name      1
track_name      1
popularity      0
duration_ms     0
explicit        0
danceability    0
energy          0
key            0
loudness        0
mode           0
speechiness     0
acousticness    0
instrumentalness 0
liveness        0
valence         0
tempo          0
time_signature  0
track_genre     0
dtype: int64
```

In [6]:

```
df.drop(['Unnamed: 0', 'track_id', 'artists', 'album_name', 'track_name'], axis=1, inplace = True)
```

In [7]:

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 114000 entries, 0 to 113999
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   popularity             114000 non-null  int64
1   duration_ms            114000 non-null  int64
2   explicit               114000 non-null  bool
3   danceability           114000 non-null  float64
4   energy                 114000 non-null  float64
5   key                    114000 non-null  int64
6   loudness               114000 non-null  float64
7   mode                   114000 non-null  int64
8   speechiness            114000 non-null  float64
9   acousticness           114000 non-null  float64
10  instrumentalness        114000 non-null  float64
11  liveness               114000 non-null  float64
12  valence                114000 non-null  float64
13  tempo                  114000 non-null  float64
14  time_signature         114000 non-null  int64
15  track_genre            114000 non-null  object
dtypes: bool(1), float64(9), int64(5), object(1)
memory usage: 13.2+ MB
```

In [8]:

df.shape

Out[8]:

(114000, 16)

In [9]:

```
df.song_duration_ms= df.duration_ms.astype(float)
df.time_signature= df.time_signature.astype(float)
```

In [10]:

df.describe()

Out[10]:

	popularity	duration_ms	danceability	energy	key	loudness	mode	speechiness	acousticness
count	114000.000000	1.140000e+05	114000.000000	114000.000000	114000.000000	114000.000000	114000.000000	114000.000000	114000.000000
mean	33.238535	2.280292e+05	0.566800	0.641383	5.309140	-8.258960	0.637553	0.084652	0.314910
std	22.305078	1.072977e+05	0.173542	0.251529	3.559987	5.029337	0.480709	0.105732	0.332523
min	0.000000	0.000000e+00	0.000000	0.000000	0.000000	-49.531000	0.000000	0.000000	0.000000
25%	17.000000	1.740660e+05	0.456000	0.472000	2.000000	-10.013000	0.000000	0.035900	0.016900
50%	35.000000	2.129060e+05	0.580000	0.685000	5.000000	-7.004000	1.000000	0.048900	0.169000
75%	50.000000	2.615060e+05	0.695000	0.854000	8.000000	-5.003000	1.000000	0.084500	0.598000
max	100.000000	5.237295e+06	0.985000	1.000000	11.000000	4.532000	1.000000	0.965000	0.996000

In [11]:

```
df["popularitybinary"] = [ 1 if i >= 50 else 0 for i in df.popularity ]
df["popularitybinary"].value_counts()
```


Out[11]:

```
0    84633
1    29367
Name: popularitybinary, dtype: int64
```

If 'probability' > 50 we labeled it "1" and if it is < 50 we labeled it "0". In this way we have "1" for the popular songs and "0" for the unpopular ones in the column df["popularitybinary"].

In [12]:

```
f,ax = plt.subplots(figsize=(13, 13))
mask = np.zeros_like(df.corr())
mask[np.triu_indices_from(mask)] = True
sns.heatmap(df.corr(), annot=True, linewidths=0.4, linecolor="white", fmt= '.1f', ax=ax, cmap="Reds", mask=mask)
plt.show()
```



A correlation heatmap showing a 2 dimensional correlation-matrix, using colored cells to represent data from usually a monochromatic scale.

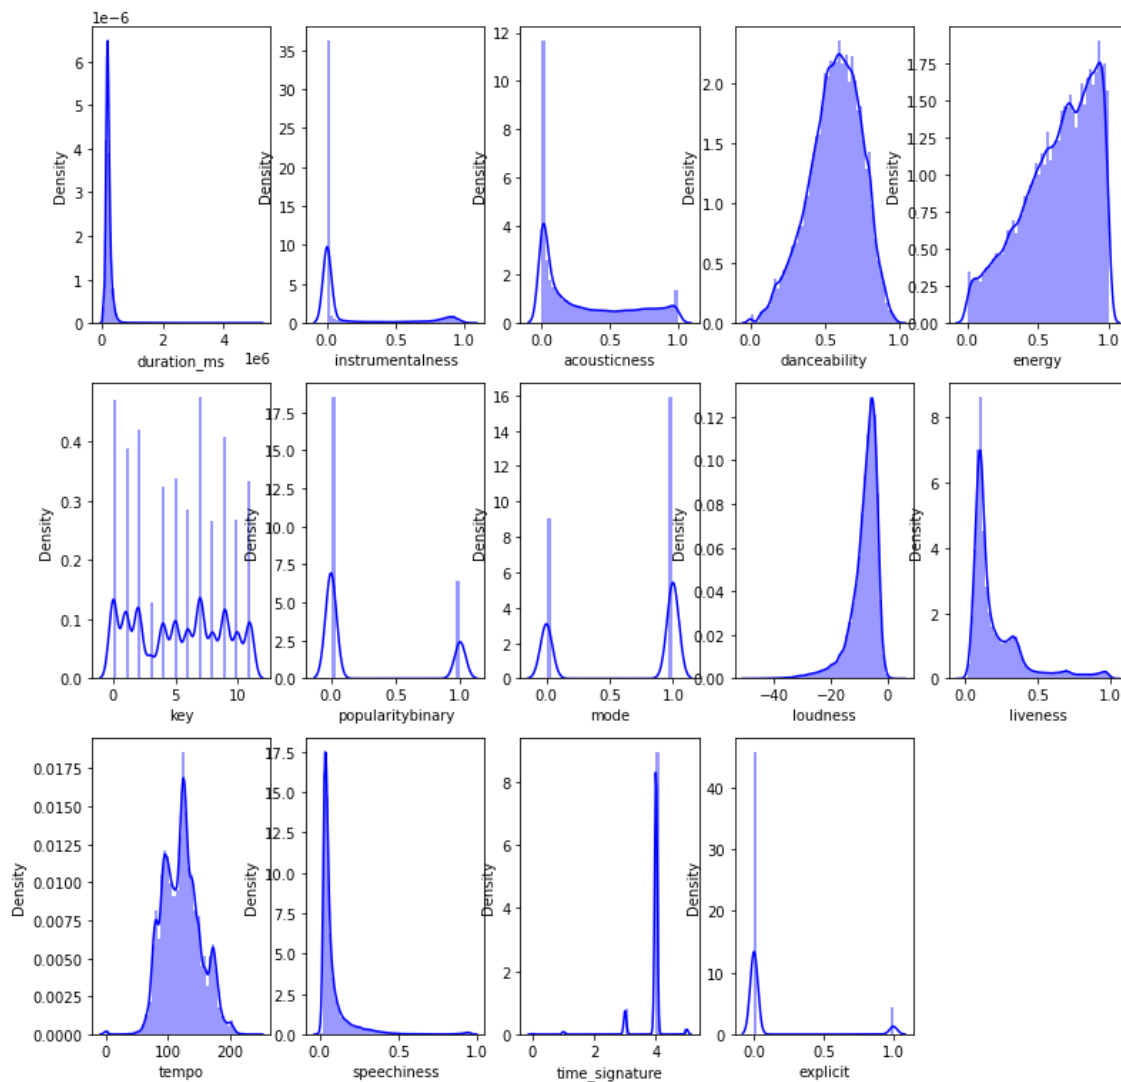
The correlation between 'loudness' and 'energy' is 0.8 which is strong. Except that all the correlations are quite low. When we compare the correlation between popularitybinary and all other features, we don't see a strong correlation (a linear relationship) that gives us a clear information about popularity. Explicit, danceability and loudness seems to have correlation with popularity feature(0.10) and istrumentalness, acousticness and liveness has (-0.10).

Features distribution

In [13]:

```
f, axes = plt.subplots(3, 5, figsize=(13, 13))
sns.distplot(df["duration_ms"], color="Blue", ax=axes[0,0])
sns.distplot(df["instrumentalness"], color="Blue", ax=axes[0,1])
sns.distplot(df["acousticness"], color="Blue", ax=axes[0,2])
sns.distplot(df["danceability"], color="Blue", ax=axes[0,3])
sns.distplot(df["energy"], color="Blue", ax=axes[0,4])
sns.distplot(df["key"], color="Blue", ax=axes[1,0])
sns.distplot(df["popularitybinary"], color="Blue", ax=axes[1,1])
sns.distplot(df["mode"], color="Blue", ax=axes[1,2])
sns.distplot(df["loudness"], color="Blue", ax=axes[1,3])
sns.distplot(df["liveness"], color="Blue", ax=axes[1,4])
sns.distplot(df["tempo"], color="Blue", ax=axes[2,0])
sns.distplot(df["speechiness"], color="Blue", ax=axes[2,1])
sns.distplot(df["time_signature"], color="Blue", ax=axes[2,2])
sns.distplot(df["explicit"], color="Blue", ax=axes[2,3])
sns.distplot(df["valence"], color="Blue", ax=axes[2,4])

f.delaxes(axes[2][4])
plt.show()
```



The distribution of songs features like danceability, energy, loudness and tempo are quite high. People like fast and loud music.

In [14]:

```
df.drop(["popularity", "track_genre", "explicit"], axis=1, inplace=True)
```

In [15]:

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 114000 entries, 0 to 113999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   duration_ms            114000 non-null  int64
1   danceability            114000 non-null  float64
2   energy                  114000 non-null  float64
3   key                     114000 non-null  int64
4   loudness                114000 non-null  float64
5   mode                    114000 non-null  int64
6   speechiness             114000 non-null  float64
7   acousticness            114000 non-null  float64
8   instrumentalness         114000 non-null  float64
9   liveness                114000 non-null  float64
10  valence                  114000 non-null  float64
11  tempo                    114000 non-null  float64
12  time_signature           114000 non-null  float64
13  popularitybinary         114000 non-null  int64
dtypes: float64(10), int64(4)
memory usage: 12.2 MB
```

In [16]:

df.columns[df.isnull().any()]

Out[16]:

Index([], dtype='object')

In [17]:

def change_type(var): df[var] = df[var].astype(int)

Data Preparation

In [18]:

```
y = df["popularitybinary"].values
#x_data=df.drop(["popularitybinary"],axis=1)

#normalization
z = (df - np.min(df))/(np.max(df)-np.min(df)).values
```

In [19]:

```
X_train = df.copy()
y_train = df.pop("popularitybinary")
```

In [20]:

```
from imblearn.over_sampling import RandomOverSampler, SMOTE
from collections import Counter
```

In [21]:

```
#to balance the distribution of popularity binary
ros = RandomOverSampler()

X_train, y_train = ros.fit_resample(X_train, y_train)

print(Counter(y_train))

Counter({1: 84633, 0: 84633})
```

In [22]:

```
print(X_train.value_counts())
```

duration_ms	danceability	energy	key	loudness	mode	speechiness	acousticness	instrumentalness	live
ness	valence	tempo	time_signature	popularity	binary				
162897	0.647	0.876	10	-5.662	1	0.1850	0.881000	0.000036	0.26
00	0.9490	151.925	4.0	0		146			
118840	0.602	0.553	11	-9.336	1	0.0328	0.108000	0.000000	0.05
12	0.9710	130.594	4.0	0		76			
172342	0.795	0.565	3	-4.457	0	0.0948	0.131000	0.000000	0.08
02	0.5500	87.925	4.0	0		73			
131733	0.579	0.502	8	-7.570	1	0.0513	0.733000	0.000000	0.28
10	0.8360	76.783	4.0	0		69			
243057	0.503	0.582	0	-4.324	1	0.0253	0.472000	0.000000	0.10
30	0.3260	77.321	4.0	0		66			
...									
198680	0.514	0.981	10	-2.063	1	0.0454	0.026100	0.000000	0.26
00	0.9230	162.844	4.0	0		1			
	0.551	0.827	10	-6.358	0	0.0600	0.204000	0.000001	0.31
80	0.6760	101.442	4.0	0		1			
	0.667	0.734	5	-7.038	0	0.0560	0.082700	0.002100	0.11
20	0.3450	92.997	4.0	0		1			
198688	0.673	0.931	10	-4.326	0	0.0394	0.000368	0.487000	0.04
35	0.9310	122.023	4.0	0		1			
5237295	0.695	0.736	5	-11.371	0	0.0374	0.003990	0.860000	0.09
10	0.0509	124.001	4.0	0		1			

Length: 84130, dtype: int64

In [23]:

```
X_train.to_csv("training.csv")
```